

Data Wrangling II

Ayush Jain

1-25-2022

Main Ideas

- To answer questions with data, we often need to use related data from many different datasets.
- We can combine data from different sources using a well-chosen join function.

Coming Up

- Homework #01 due Thursday.
- Lab #03 due Friday

Lecture Notes and Exercises

```
library(tidyverse)
```

Instead of working with a single dataset, usually you will have to work with many different related datasets. To answer research questions using related datasets, we need to develop tools to join datasets together.

There are many possible types of joins. All have the format `something_join(x, y)`.

- `inner_join()`: join all rows from `x` where there are matching values in `y`. Return all combinations in case of multiple matches
- `left_join()`: include all rows from `x`
- `right_join()`: include all rows from `y`
- `full_join()`: include all rows in `x` or `y`
- `semi_join()`: return all rows from `x` with match in `y`
- `anti_join()`: return all rows from `x` without a match in `y`

```
x <- tibble(value = c(1, 2, 3),
             xcol = c("x1", "x2", "x3"))
y <- tibble(value = c(1, 2, 4),
             ycol = c("y1", "y2", "y4"))
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
```

```
## 1      1 x1
## 2      2 x2
## 3      3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1      1 y1
## 2      2 y2
## 3      4 y4
```

We will demonstrate each of the joins on these small, toy datasets.

```
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
## 1      1 x1
## 2      2 x2
## 3      3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1      1 y1
## 2      2 y2
## 3      4 y4
```

```
inner_join(x, y)
```

```
## Joining, by = "value"
```

```
## # A tibble: 2 x 3
##   value xcol ycol
##   <dbl> <chr> <chr>
## 1      1 x1   y1
## 2      2 x2   y2
```

```
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
## 1      1 x1
## 2      2 x2
## 3      3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1     1 y1
## 2     2 y2
## 3     4 y4
```

```
left_join(x, y)
```

```
## Joining, by = "value"
```

```
## # A tibble: 3 x 3
##   value xcol ycol
##   <dbl> <chr> <chr>
## 1     1 x1   y1
## 2     2 x2   y2
## 3     3 x3   <NA>
```

```
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
## 3     3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1     1 y1
## 2     2 y2
## 3     4 y4
```

```
right_join(x, y)
```

```
## Joining, by = "value"
```

```
## # A tibble: 3 x 3
##   value xcol ycol
##   <dbl> <chr> <chr>
## 1     1 x1   y1
## 2     2 x2   y2
## 3     4 <NA> y4
```

```
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
## 3     3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1     1 y1
## 2     2 y2
## 3     4 y4
```

```
full_join(x, y)
```

```
## Joining, by = "value"
```

```
## # A tibble: 4 x 3
##   value xcol ycol
##   <dbl> <chr> <chr>
## 1     1 x1 y1
## 2     2 x2 y2
## 3     3 x3 <NA>
## 4     4 <NA> y4
```

```
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
## 3     3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1     1 y1
## 2     2 y2
## 3     4 y4
```

```
semi_join(x, y)
```

```
## Joining, by = "value"
```

```
## # A tibble: 2 x 2
##   value xcol
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
```

```
x
```

```
## # A tibble: 3 x 2
##   value xcol
##   <dbl> <chr>
## 1     1 x1
## 2     2 x2
## 3     3 x3
```

```
y
```

```
## # A tibble: 3 x 2
##   value ycol
##   <dbl> <chr>
## 1     1 y1
## 2     2 y2
## 3     4 y4
```

```
anti_join(x, y)
```

```
## Joining, by = "value"
```

```
## # A tibble: 1 x 2
##   value xcol
##   <dbl> <chr>
## 1     3 x3
```

How do the join functions above know to join `x` and `y` by `value`? Examine the names to find out.

```
names(x)
```

```
## [1] "value" "xcol"
```

```
names(y)
```

```
## [1] "value" "ycol"
```

We will again work with data from the `nycflights13` package.

```
library(nycflights13)
```

Examine the documentation for the datasets `airports`, `flights`, and `planes`.

Question: How are these datasets related? Suppose you wanted to make a map of the route of every flight. What variables would you need from which datasets? Answer: From `airports`, you would need name, lat, lon. From `flights`, you would need origin, dest, and distance.

Join `flights` to `airports`. Note these two datasets have no variables in common so we will have to specify the variable to join by using `by =`. Check out the documentation for more information.

```
flights %>%
  left_join(airports, by = c("dest" = "faa"))

## # A tibble: 336,776 x 26
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2     830           819
## 2  2013     1     1     533           529           4     850           830
## 3  2013     1     1     542           540           2     923           850
## 4  2013     1     1     544           545          -1    1004          1022
## 5  2013     1     1     554           600          -6     812           837
## 6  2013     1     1     554           558          -4     740           728
## 7  2013     1     1     555           600          -5     913           854
## 8  2013     1     1     557           600          -3     709           723
## 9  2013     1     1     557           600          -3     838           846
##10  2013     1     1     558           600          -2     753           745
## # ... with 336,766 more rows, and 18 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   name <chr>, lat <dbl>, lon <dbl>, alt <dbl>, tz <dbl>, dst <chr>,
## #   tzone <chr>
```

Practice

- (1) Create a new dataset `dest_delays` with the median arrival delay for each destination. Note this question does not require you to use joins.

```
dest_delays <-
  flights %>%
  group_by(dest) %>%
  summarise(median_arr_delay = median(arr_delay, na.rm= TRUE))
```

- (2) Create a new dataset by joining `dest_delays` and `airports`. Only include observations that have both delay and airport information. Note `dest_delays` and `flights` have no variables in common so you will need to specify the variables to join using `by` as in the example above.

```
dest_delays %>%
  inner_join(airports, by = c("dest" = "faa"))

## # A tibble: 101 x 9
##   dest median_arr_delay name          lat   lon alt   tz dst tzone
##   <chr>              <dbl> <chr>         <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 ABQ                -5.5 Albuquerque I~ 35.0 -107.  5355   -7 A   America~
```

```
## 2 ACK -3 Nantucket Mem 41.3 -70.1 48 -5 A America~
## 3 ALB -4 Albany Intl 42.7 -73.8 285 -5 A America~
## 4 ANC 1.5 Ted Stevens A~ 61.2 -150. 152 -9 A America~
## 5 ATL -1 Hartsfield Ja~ 33.6 -84.4 1026 -5 A America~
## 6 AUS -5 Austin Bergst~ 30.2 -97.7 542 -6 A America~
## 7 AVL -1 Asheville Reg~ 35.4 -82.5 2165 -5 A America~
## 8 BDL -10 Bradley Intl 41.9 -72.7 173 -5 A America~
## 9 BGR -9 Bangor Intl 44.8 -68.8 192 -5 A America~
## 10 BHM -2 Birmingham In~ 33.6 -86.8 644 -6 A America~
## # ... with 91 more rows
```

Question: Are all of the variables in `dest_delays` included in the new dataset you created by joining `dest_delays` and `airports`? Use an appropriate join function to investigate this issue and determine what is going on here. Answer: No, there are 4 rows that are in `dest_delays` but not in `airports`

Use an `anti_join` to help diagnose this issue. Recall `anti_join` returns all rows from `x` without a match in `y`, so it will return all rows in `dest_delays` that don't have a match in `airports`.

```
dest_delays %>%
  anti_join(airports, by = c("dest" = "faa"))
```

```
## # A tibble: 4 x 2
##   dest median_arr_delay
##   <chr>           <dbl>
## 1 BQN             -1
## 2 PSE              0
## 3 SJU             -6
## 4 STT             -9
```

(3) Is there a relationship between the age of a plane and its delays? The plane tail number is given in the `tailnum` variable in the `flights` dataset. The year the plane was manufactured is given in the `year` variable in the `planes` dataset.

- Step #1: Start by finding the average arrival delay for each plane and store the resulting dataset in `plane_delays`.

```
plane_delays <- flights %>%
  group_by(tailnum) %>%
  summarise(mean_arr_delay = mean(arr_delay, na.rm = TRUE))
```

- Step #2: Join `plane_delays` to the `planes` data using an appropriate join and then use `mutate` to create an `age` variable. Note this data is from 2013.

```
plane_delays %>%
  left_join(planes, by = "tailnum") %>%
  mutate(age = 2013 - year)
```

```
## # A tibble: 4,044 x 11
##   tailnum mean_arr_delay year type manufacturer model engines seats speed
##   <chr>           <dbl> <int> <chr>      <chr>      <chr>   <int> <int> <int>
## 1 D942DN          31.5     NA <NA>      <NA>      <NA>     NA    NA    NA
```

```
## 2 NOEGMQ          9.98      NA <NA>      <NA>      <NA>      NA      NA      NA
## 3 N10156         12.7      2004 Fixed w~ EMBRAER      EMB~      2      55      NA
## 4 N102UW          2.94      1998 Fixed w~ AIRBUS INDUS~ A320~      2     182      NA
## 5 N103US         -6.93      1999 Fixed w~ AIRBUS INDUS~ A320~      2     182      NA
## 6 N104UW          1.80      1999 Fixed w~ AIRBUS INDUS~ A320~      2     182      NA
## 7 N10575         20.7      2002 Fixed w~ EMBRAER      EMB~      2      55      NA
## 8 N105UW         -0.267     1999 Fixed w~ AIRBUS INDUS~ A320~      2     182      NA
## 9 N107US         -5.73      1999 Fixed w~ AIRBUS INDUS~ A320~      2     182      NA
## 10 N108UW        -1.25      1999 Fixed w~ AIRBUS INDUS~ A320~      2     182      NA
## # ... with 4,034 more rows, and 2 more variables: engine <chr>, age <dbl>
```


- Step #3: Finally, create an effective visualization of the data.

```
plane_delays %>%
  left_join(planes, by = "tailnum") %>%
  mutate(age = 2013 - year) %>%
  ggplot(aes(x = age, y = mean_arr_delay)) +
  geom_point() +
  geom_smooth() +
  theme_classic()
```

```
## 'geom_smooth()' using method = 'gam' and formula 'y ~ s(x, bs = "cs")'
```

```
## Warning: Removed 798 rows containing non-finite values (stat_smooth).
```

```
## Warning: Removed 798 rows containing missing values (geom_point).
```



Additional Resources

- <https://rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf>